

ENEEL KUMAR A

+91-9606288164 | veneel77work@gmail.com | [linkedin.com/in/veneelkumara77](https://www.linkedin.com/in/veneelkumara77) | github.com/Veneel77 | Bangalore, India

SUMMARY

Data Scientist with hands-on experience building end-to-end ML and statistical modeling pipelines, AI-powered applications, and data analytics solutions using Python, SQL, and open-source frameworks. Skilled in problem scoping, feature engineering, model selection, and translating analytical results into actionable insights for stakeholders. Published researcher in deep learning with a strong foundation in data structures, version control, and software engineering practices. Azure AI-900 certified; experienced across AWS, Azure, and GCP.

SKILLS

Data Science & Statistics: Statistical Analysis, Predictive Modelling, Feature Engineering, Model Evaluation, Hypothesis Testing, A/B Testing Concepts

Programming: Python (primary), SQL, R (familiar), Pandas, NumPy, PySpark

ML / AI: Scikit-Learn, XGBoost, TensorFlow, Keras, LLMs (GPT, LLaMA), LangChain, CrewAI, RAG Architecture

Databases & Querying: MySQL, PostgreSQL, MongoDB, ChromaDB (Vector DB), SQL & NoSQL querying

Data Engineering: Apache Spark, Airflow, ETL/ELT Pipelines, Data Processing

Visualisation & Reporting: Power BI, Streamlit, Jupyter Notebook

Cloud Platforms: AWS (SageMaker, S3, EC2, IAM), Azure (AI Services, Functions, App Services), GCP

Software Engineering: Git, CI/CD, Docker, REST APIs, FastAPI, SDLC, OOP, Code Review

EXPERIENCE

AI/ML Engineer Trainee

Sep 2024 – Feb 2025

Rooman Technologies · Remote · 6 months

- Developed and evaluated Python-based ML models (classification, regression, ensemble methods) using Scikit-Learn; applied structured feature engineering, model benchmarking, and performance tuning to optimise predictive accuracy.
- Built a Heart Disease Prediction system using ensemble methods (Random Forest, Gradient Boosting) — implemented end-to-end pipeline from data preprocessing to evaluation; documented findings and presented model results to project stakeholders.
- Applied SDLC best practices including code review, debugging, and technical documentation throughout the development lifecycle.

PROJECTS

Customer Churn Prediction System

Sep 2025

Python · XGBoost · Scikit-Learn · Pandas · NumPy · Power BI | github.com/Veneel77/Customer-Churn-Prediction-System

- Built a complete data science pipeline on the IBM Telco dataset (7,000+ records, 21 features) — data cleaning, exploratory analysis, feature engineering, and benchmarking of multiple models; XGBoost achieved ~85% accuracy.
- Translated model results into a stakeholder-facing Power BI dashboard, surfacing churn drivers and segment-level insights to support business decision-making.
- Documented the full modelling methodology and analytical conclusions as a structured technical report.

RAG-Based Knowledge Assistant

Nov – Dec 2025

Python · LangChain · ChromaDB · HuggingFace · FastAPI · Streamlit | github.com/Veneel77/RAG-based-Knowledge-assistant

- Designed and built an end-to-end document-aware question-answering system using retrieval-augmented generation (RAG); demonstrated ability to scope a proof-of-concept and build a modular, scalable solution.
- Integrated HuggingFace sentence-transformer embeddings with ChromaDB for semantic retrieval; implemented unit tests for ingestion and retrieval pipelines; open-sourced under MIT licence with full documentation.

Agentic Data Pipeline Reliability Platform

Jun 2026 – Present

Python · CrewAI · Google Gemini 2.5 Flash · ChromaDB · FastAPI · Docker | github.com/Veneel77

- Developed a 5-agent AI orchestration system automating ETL incident lifecycle management — failure classification, root cause analysis, fix recommendation, and report generation; architected for scale and reproducibility.
- Built RAG-powered semantic search over Spark/Airflow knowledge bases; productionised via a FastAPI backend and Docker Compose deployment with structured logging and observability hooks.

RESEARCH & PUBLICATIONS

Alzheimer's Diagnosis via Deep Learning

JETIR, 2025

- Published peer-reviewed research applying CNN architectures and transfer learning (VGG16, ResNet50) on MRI datasets for multi-class Alzheimer's stage classification; achieved >90% accuracy through systematic hyperparameter tuning and data augmentation. Journal of Emerging Technologies and Innovative Research (JETIR), 2025.

EDUCATION

KSSEM, Bangalore — B.E. (Engineering) in AI & Data Science

Aug 2021 – Jun 2025

Relevant coursework: Machine Learning, Deep Learning, Statistics, Cloud Computing, NLP, Data Mining & Warehousing, Database Management Systems

CERTIFICATIONS

- Microsoft Azure AI-900
- Oracle GenAI Professional
- Google Cloud Kubernetes
- Google Advanced Data Analytics

HONORS & AWARDS

- 1st Prize — AI Hackathon (College Level)
- 210+ Solved LeetCode Problems (DSA)
- Published Research Paper — JETIR 2025